




General Liquidity

Pioneering agentic vibe trading.

Tiberiu Toca

Founder, CEO

contact@generalliquidity.com

The thesis — worldbuilding

The economy is more programmable, more adversarial, and harsher at scale than at any point in modern history. Financial institutions already operate production-grade AI: quant systems, state-of-the-art economists, and microstructure models, positioned so whales can beat retail positions by design. AI has killed manual trading at the institutional end. Agents are eating the market. In parallel, the career ladder for young people is losing its bottom rungs. Cognitive entry-level work was the training ground for every professional career of the last century, and cognitive entry-level work is exactly what agents do best. What analysts call the hollowing out of the corporate structure is the same phenomenon this generation is living through as *"there are no jobs and the ones that exist don't lead anywhere."*

There is a wealth transfer happening right now, caused by AI itself, and the only question is whether ordinary people get to participate or get to watch. Research from Allianz already shows that software investment produces the strongest labor-substitution effect of any capital category, a **0.22–0.29% labor reduction per 1% increase in software investment**, with finance and real estate facing the steepest employment declines. AI agents are a new factor of production, neither classical labor nor classical capital, what the systemic analyses are starting to call **agentic capital** or **digital labor**. The returns to that factor flow to whoever owns it. Right now that is the platform layer and the institutional trading floor. It does not have to stay that way.

The peer-reviewed literature is catching up to what the institutional floor already knows. The Google DeepMind *Virtual Agent Economies* paper (arxiv 2509.10147, 2025) names the mechanism precisely: *"More capable AI assistants negotiate better deals for their users."* The implication is not subtle. In every market where AI agents transact on their principal's behalf, the quality of the agent determines the deal. A retail investor with a 2016-era chatbot recommendation is not competing with a hedge fund's full-stack execution infrastructure. They are being extracted from. The paper further documents the systemic risks: flash-crash dynamics from correlated agentic behavior, and the critical finding that multi-tiered oversight frameworks, exactly the kind Gordon's six-mode permission model implements, are the primary known mitigation for agentic market instability. The academic literature and Gordon's architecture are pointing at the same design.

The participation gap is not about access. Most people who will benefit from Gordon already have a brokerage account and a crypto wallet. They are inside the system. The problem is that they are outside the loop that matters: the loop where a thesis becomes a plan, a plan becomes a position, a position is monitored, and outcomes compound into judgment. Having an account is not the same as operating an account. Participation is not agency.

General Liquidity exists to make sure the agentic generation gets to participate, not just watch. We are building **agentic vibe trading**: a consumer AI product category in which a person with no quant background, no algo experience, no economist pedigree, and no finance degree can operate in markets with the same control plane institutions already use. Trading is the wedge because it is where wealth gets extracted from markets today. After that we extend into personal finance, and the context each operator accumulates (theses, trades, preferences, discipline, outcomes) compounds into a **personal financial mobility engine**.

Programmability without agency is a bigger trap than no programmability at all. More rails, more models, more access, and a population that understands less and less of what is happening to their money. The field's most common responses to this problem all collapse toward the wrong answer. Foundation models absorbing the application layer, they commoditize the model, they do not own the operator context built on top of it. New rails making legacy infrastructure obsolete, DeFi and TradFi will coexist for decades; the productive position is the seam between them, not a bet on migration. Natural language solving the software problem, natural language is an interface, not a control plane; a chatbot that reasons aloud but cannot execute is a more persuasive version of the same access problem, not an answer to it. In each case the surface changes. The asymmetry remains.

The control plane between raw capability and operator intent is the product. **Access is not agency. Raw rails do not create trust. Raw models do not create judgment.** That plane is what General Liquidity builds.

The opportunity — the agentic economy is happening now

AI is projected to contribute \$15 trillion to the global economy by 2030. Agentic AI specifically is projected to add \$2.6–4.4 trillion globally over the next 2–5 years (Allianz). Those are macro numbers. The financial services layer has its own.

The CFA Institute confirms that 73% of Y Combinator investments between January 2024 and June 2025 were in agentic AI companies. Earnings call mentions of AI agents increased 4× quarter-over-quarter in Q4 2024 alone. Capital and corporate attention are already moving; the infrastructure buildout is happening in real time. The agentic AI market overall is moving from \$4.54B in 2025 to a projected \$98.26B by 2033 at a 46.87% CAGR (DataM Intelligence, February 2026). Within that, the retail algorithmic trading segment alone is projected to grow from \$3.55B in 2024 to \$7.18B by 2030 (Grand View Research). A 46.87% CAGR is not a category trend. It is a generational platform shift.

Pool	Scale	Source
Global payments industry annual revenue	\$2.7T	McKinsey
Consumer credit-card revenue (2024)	\$234B	McKinsey
Global retail-bank net interest income share of profit	~30%	McKinsey
Addressable autopilot-services market (aggregate)	\$1.5T+	Sequoia
Services budget per dollar of software budget	6:1	Sequoia

Sequoia's autopilot opportunity map breaks the services number into specific verticals, each its own trillion-dollar-class pattern.

Vertical	Labor TAM	Current autopilot leaders
Insurance brokerage	\$140–200B	WithCoverage, Harper
Claims adjusting	\$50–80B	Pace, Strala
Healthcare revenue-cycle	\$50–80B	Anterior
Accounting & audit	\$50–80B	Rillet, Basis
Tax advisory	\$30–35B	TaxGPT, Skalar, Ravical
Legal transactional	\$20–25B	Harvey, Crosby, Lawhive
IT managed services	\$100B+	Edra, Serval
Supply chain / procurement	\$200B+	Magentic, AskLio, Tacto
Recruitment & staffing	\$200B+	Juicebox, Mercor, Jack & Jill
Management consulting	\$300–400B	early — TBD

Gordon's wedge, delegated trading for retail operators is the consumer-finance analog of every line on that table. Intelligence-heavy, currently outsourced (to brokers, advisors, influencers, or sheer luck), and waiting for the first credible autopilot to collapse the cost of doing it well.

Adoption is no longer theoretical on either side. Institutional, retail, and regulatory signals are converging inside a 12-month window:

- **Gartner:** 33% of enterprise software will include agentic AI by 2028, up from under 1% in 2024. 15% of day-to-day work decisions autonomous by 2028, from 0%. 80% of standard customer-service queries by 2029.
- **Wolters Kluwer / AImagicx:** 44% of finance teams are already using agentic AI in Q1 2026, a 600% year-over-year increase. The adoption curve in financial services specifically is vertical.
- **Deloitte:** 50% of companies that have implemented generative AI will deploy agentic AI pilots by 2027. The floor is rising rapidly.
- **Cleo AI** (February 5, 2026): Launched "Autopilot," a multi-agent architecture that takes financial action within user-set guardrails. Their research found 53% of Americans say they would trust AI to advise on disposable income and 52% are open to AI automatically moving money to avoid overdrafts. Consumer readiness is not theoretical.
- **PagerDuty:** 51% of companies already have agents in production. 78% plan adoption. 90% of non-tech firms are on board.
- **Goldman Sachs:** 300M FTE jobs globally exposed to generative-AI automation.
- **Yu'e Bao** (China, money-market fund app) moved \$268B of consumer cash in under four years once the rate gap became visible. Consumers do move money when they can see the opportunity.
- **Raisin** (Europe, deposit marketplace) has \$80B across 250 banks. Same thesis, different geography.
- **Santander + Mastercard** completed Europe's first live end-to-end AI-executed regulated financial transaction in March 2026. Not a pilot. A live payment with delegated authority.
- **DTCC** reports its Risk Calculator AI Assistant runs at ~97% accuracy with zero hallucinations, cutting client-meeting prep from one week to one day and driving a fourfold productivity lift for client-facing teams. This is the plumbing of global capital markets publicly admitting agents already work at scale.
- **Moody's** reports its AI Research Assistant users consume 60% more research while cutting task completion times by 30%. Over 90% of AI interactions are now focused on high-value analytics.
- **Coinbase** is publicly pivoting to become mission-critical infrastructure for the Agentic Web: AgentKit, x402 with Cloudflare, World ID for Sybil resistance, Base + USDC as agent settlement rails. Wall Street is pricing the pivot in: consensus EPS of \$3.79 (2026) growing to ~\$7.11 (2028), forward P/E moving from 42x to 24x on that trajectory. Coinbase's pivot makes structural sense: crypto increasingly functions as machine-native financial infrastructure: programmable, composable, and built in a form software can use natively, more than it functions as a human-first consumer product. A machine economy is materializing on rails that were finally designed for machines.
- **Human.tech's Agentic WaaP** (Wallet-as-a-Protocol) launched at WalletCon 2026. Its broader identity infrastructure protects \$500M+ from Sybil attacks, secured by \$3B+ in restaked ETH.
- **MIT + Stanford** field study on AI augmentation: 14% productivity lift on average across customer-service workers, 34% for less-experienced workers. AI is a great equalizer. The biggest lift goes to the people furthest from the frontier. That is the empirical foundation of the economic-mobility thesis.
- Operationally, the plumbing is ready. FedNow (US), TIPS and SEPA Instant (EU), and UPI (India) have eliminated settlement lag. Stablecoin rails (Visa-USDC, Circle, Coinbase Payments) add programmable, sub-second transfers. The rails an agent needs to sweep funds, split orders, or pay a vendor in real time are live today. The bottleneck is the control plane on top, not the plumbing underneath.

The consumer side is still mostly empty. The institutional rails are shipping; the retail control plane is not. That is the gap.

The paradigm shift we ride

Four authoritative voices have published the same thesis inside the last 18 months, each pointing at the exact wedge General Liquidity is building. Sequoia is the most applicable to our go-to-market. Block is the closest to our architecture. Pear VC's Request for Startups names three categories we hit directly. McKinsey has put the most detailed version on the institutional side. And the academic literature has a paper with the exact title we would have written if we had written it. When that many credible voices converge on the same claim in the same year, it is a market forming.

Sequoia — *Services, the New Software* (Julien Bek, March 2026). *"If you sell the tool, you're in a race against the model. But if you sell the work, every improvement in the model makes your service faster, cheaper, and harder to compete with."* Copilots to autopilots. Sequoia's specific wedge-strategy argument is directly applicable to Gordon's go-to-market: *"Replacing an outsourcing contract with an AI-native services provider is a vendor swap. Replacing headcount is a reorg."* Gen Z and younger millennials already outsource financial judgment: to influencers, to TikTok, to Robinhood's default UX, to sheer guesswork. Replacing that informal outsourcing with Gordon is a clean vendor swap. We do not need to convince anyone to fire a human advisor. We need to be measurably better than the implicit advisor they already rely on. Sequoia also draws the line between intelligence and judgment: *"Writing code is mostly intelligence. Knowing what to build next is judgement."* The same split defines trading. Executing trades: scanning, routing, sizing, timing is intelligence work, and agents do it better, faster, and cheaper. Knowing which thesis to deploy, in what market regime, with what conviction, at what sizing; that is judgment work, and it still belongs to the operator. Gordon's architecture is this split made concrete. *"Today, more tasks are started by agents than by humans,"* Sequoia writes about Cursor. The same inflection is arriving in retail finance. *"The work budget in any profession dwarfs the tool budget, and autopilots capture the work budget from day one."*

Block — *From Hierarchy to Intelligence* (Jack Dorsey and Roelof Botha, March 2026). *"A world model that can't touch the world is just a database."* The four-layer model: capabilities, world model, intelligence layer, interfaces is the pattern Gordon is built around. We are building Block's "cognitive enterprise" pattern as a consumer product. Block also articulates an operational principle we have adopted: *"customer reality generates the backlog directly."* When an agent tries to compose a solution and cannot because a capability does not exist, that failure signal is the next build. We do not hypothesize a roadmap. We observe the roadmap emerging from real operator sessions. Block further makes the most direct argument for why financial behavior data is uniquely valuable: *"Money is the most honest signal in the world. People lie on surveys. They ignore ads. They abandon carts. But when they spend, save, send, borrow, or repay, that's the truth."* Gordon's per-operator financial context is built on this signal. Not aggregate behavioral proxies. Not ad-targeting data. The operator's own revealed preferences, compounding under the pressure of real capital over real time.

QED Investors — *The OpenClaw Phase* (Amias Gerety, April 2026). QED is among the most credible fintech-focused venture firms globally, backers of Credit Karma, Nubank, and Remitly. Partner Amias Gerety declared in April 2026 that financial services is transitioning from AI copilot to the "OpenClaw" phase: autonomous reasoning agents that *"complete entire workflows independently."* This is not a general AI thesis. It is a fintech-specific declaration from a firm with direct sight lines into what the best-funded consumer finance companies are building. Gordon is designed for precisely this phase: not to recommend, but to complete.

Pear VC — *Request for Startups* (S26 cohort). Three of Pear's twelve RFS categories are directly Gordon-shaped, and the primary match is the first one Pear lists. **Category 1, Plan mode for every knowledge worker:** Gordon has a literal plan permission mode and a plan-first operator workflow (plan, preview, approve, execute, monitor, reconcile) that brings structured planning to financial decision-making. This is the exact pattern Pear describes for sales, marketing, and finance, except the domain is trading. Pear's framing for why this matters: current AI tools produce outputs that are "60–70% of the way there", enough to generate the analysis, not enough to close the trade. A plan→preview→approve workflow closes the gap by keeping human judgment at the decision point, not just the research stage. **Category 2, Financial-grade agent infrastructure** : trust, policy, and execution layers for agents operating in regulated finance. **Category 8, Intent-native commerce:** agents as first-class buyers, programmatic pricing, catalogs as APIs. We hit all three.

McKinsey — *The End of Inertia* (August 2025). *"Margin accrues in the gap between what consumers could do and what they actually do. Agentic AI breaks that equation."* McKinsey names five control points in the agentic economy: credentialing and identity, trust and liability wrappers, merchant and platform integration, decisioning logic, and behavioral data with intent signals. Gordon owns the operator-facing surface of all five. McKinsey adds a competitive warning: *"Relevance must be re-earned moment by moment, every time an agent refreshes its rankings."* Platform-aligned agents that depend on brand inertia will fail this test. Gordon is designed to pass it by construction.

The Agentic Economy (arxiv 2505.15799, academic spine of the field). The paper argues the central economic shift is from the attention economy to the **preference economy**. *"Success will hinge on attracting early, engaged users who provide valuable feedback."* Gordon does not sell attention. We have no ad inventory. We sell operator preference data that compounds into a durable personal financial context. The flywheel works on subscriptions alone, and the data moat here is a fundamentally different asset class from the one consumer social products built.

The operational vocabulary converging across these sources is **human above the loop**, not *human in the loop* (Pawel Gmyrek, ILO, via WEF 2024). *In the loop* implies the human is a step in the pipeline. *Above the loop* implies the human sets the objective and supervises. Gordon's six-mode permission model is the cleanest implementation of human-above-the-loop currently shipping in consumer finance.

The problem — the agency bottleneck

Retail finance has been sold a story that access equals empowerment. Commission-free brokerages, neobanks, DeFi wallets, crypto exchanges, all expanded access, none created agency. A person with a Robinhood account and a Coinbase account participates in markets without operating them. They can click; they cannot reason. When prices move, they react. The tools tell them *what* is happening and leave the *what now* entirely on the user. A user who usually has no finance background, no economics training, no access to institutional research, and no infrastructure to turn a market view into a disciplined position.

Institutions, meanwhile, run fully systematized operations: research, thesis, sizing, execution, risk, review, compounding. Renaissance Technologies, Jane Street, HRT, Citadel, the top tier has iterated on systematic execution for decades. An institutional trader with a \$200k salary has agents managing positions while a 24-year-old with a \$40k salary is squinting at a candlestick chart on their phone. *The asymmetry is the product financial institutions sell.* SP Global's analysis of market fragmentation confirms that the proliferation of venues, more exchanges, more dark pools, more retail platforms has not helped retail traders navigate those markets. It has given institutional arbitrageurs more surfaces to extract from. The numbers make the gap precise: 60–75% of all US equity volume is already algorithmic (Tabb Group, 2024). Only 7–13% of human retail traders achieve positive P&L on prediction markets, versus 37%+ for AI agents operating on the same markets (Olas/LayerHub, 2026). The gap is not closing. It is widening with every improvement in institutional AI.

McKinsey documents the specific cost of inertia: US banks paid 0.07% on checking accounts and 0.38% on savings accounts in 2024, while online savings rates exceeded 4%. Consumers absorb more than three percentage points of lost yield every year simply because switching is friction-heavy enough to be avoided. The Yu'e Bao and Raisin evidence shows what happens when friction is removed: \$268B and \$80B moved, respectively, once the gap became visible and acting on it required no effort. Gordon's ask mode and plan-first execution is the architectural answer: the gap between what the operator should do and what they actually do is closed not by financial education but by reducing the cost of acting on what they already know.

What is changing is that the tools necessary to close this asymmetry finally exist, and they create a governance problem the industry is only starting to name. Hanah-Marie Darley (Geordie AI), speaking on Santander and Mastercard's March 2026 milestone: *"Permissions and governance must travel with the agent: not sit in static policy documents. The rules that define what the agent is allowed to do need to travel with it as it moves across tools, services, and markets."* Every retail fintech shipping AI today is failing this test. LLM chatbots that *recommend* without executing. Bolt-on credentials that do not carry policy context. Gordon's six-mode permission truth table, enforced at every stage of a decision, is an architectural answer to Darley's frame, built before the milestone forced the industry to articulate it.

The academic literature adds a second taxonomy that clarifies the competitive landscape. Most current AI products are *siloed* service agents: Amazon's Rufus operates only inside Amazon, Expedia's Romie operates only inside Expedia. They cannot carry context across venue boundaries or represent the operator's full picture. Gordon is an *end-to-end* agent: cross-venue, cross-market, persistent context that travels with the operator. The literature further distinguishes *unscripted* interactions (technically possible via natural language) from *unrestricted* interactions (sanctioned by governance and market structure). Natural language gets to unscripted. The permission model, the audit log, and the risk gates get to unrestricted. Most retail AI products are unscripted. Gordon is designed to be both.

A final framing from our own essays matters here. Agent commerce and by extension, agent-delegated finance splits the world into tourists who treat every transaction as an isolated one-off, and locals who operate inside supplier relationships, negotiated terms, and accumulated context. Retail investors are tourists in markets today. Gordon turns them into locals.

Agentic AI is the first technology that can close the institutional-retail asymmetry, and it is closing it on the wrong side first. OpenAI Operator, Manus AI, Perplexity, and Coinbase AgentKit are already giving sophisticated users agentic tooling. The institutional end is accelerating faster than the retail end. If nothing intervenes, the wealth transfer AI is causing will skew exactly the way every prior technological wealth transfer has skewed: toward capital, away from labor, away from the young.

The product — Gordon, and the product ladder

Gordon is the first agentic vibe trading product. It operates across crypto exchanges and stock brokers, runs a chat-first operator workflow, and carries real execution authority under a human-controlled permission model.

- **Chat-first operator workflow.** Scan, analyze, plan, preview, approve, execute, monitor, reconcile. No dashboards, no cockpit, no twelve tabs.
- **Multi-venue by design.** Crypto exchanges (Binance, Coinbase, Kraken, Bitfinex, Gemini, OKX, Hyperliquid) and stock brokers (Alpaca, Schwab, Interactive Brokers, E*TRADE, Tastytrade, TradeStation, Tradier, Trading 212, Webull), plus onchain execution across Solana, EVM, Polkadot, and Base.
- **Deep research and analysis.** Market discovery, multi-timeframe technical analysis, onchain and DEX flows, fundamentals across stocks, ETFs, indices, bonds, crypto, and macro, SEC filings, news entity extraction, plus a backtesting engine with Monte Carlo and walk-forward validation.
- **Six permission modes.** auto, ask, strict, paper, observe, plan. A single truth table that lets the operator dial autonomy from read-only research all the way to fully autonomous execution within risk gates. The same policy is enforced before, during, and at every venue call. This is the "permissions travel with the agent" pattern in shipping form.
- **Plan-first execution.** Every trade is previewed as a structured diff before anything hits a venue. Bracket, grid, DCA, and OCO plans with trailing stops and partial take-profits. No order leaves without a clean reason and an explicit approval step under the operator's mode of choice.
- **Real safety surface.** Input guardrails, access control, rate limits, risk gates, and full audit logging. A system an operator can actually trust with capital.

The product ladder

1. **Gordon (now).** Agentic vibe trading for serious operators. The Claude Code / Codex of vibe trading: terminal-native, keyboard-sovereign, scriptable, auditable. The proving ground.
2. **Gordon Desk (next).** The Cursor / Windsurf / Devin / Claude Cowork of vibe trading: a visual workspace that brings research, planning, execution, monitoring, and review into one environment. Multi-agent, multi-operator, shared context. Think *agentic Bloomberg Terminal × Cursor × Claude Projects* for markets. **Agentic multiplayer:** where shared theses, collaborative backtesting, and team trading become normal.
3. **Gordon Personal Finance OS (later, inspired by OpenClaw and Hermes).** The post-wallet, post-broker financial home. Gordon transacts on the operator's behalf across cards, banks, wallets, exchanges, payroll, tax, and subscriptions. Every operator's financial context: theses, trades, preferences, discipline, outcomes, becomes a compounding personal asset that the agent operates from. This is where the operator stops being a tourist in their own financial life and becomes a local. *"The wallet opened the door. It did not finish the room."*

A critical implication of step 3, flagged in Zuehlke's C-suite trend paper and increasingly visible in Visa and Mastercard's agent-commerce product lines: when Gordon transacts on the operator's behalf, the operator becomes Gordon's principal and every external product must court Gordon, not the human. The future customer of insurance, banking, subscriptions, groceries, and financial services is not a person but an algorithm operating on a person's mandate. Businesses that are legible, API-first, and programmatically negotiable will win. Businesses that depend on brand inertia and attention capture will lose. Gordon is designed to be the agent those mandates flow through, which makes it the entry point to a far larger market than "trading" suggests.

Gordon CLI is not a waypoint on the way to Gordon Desk. It is the test. If delegation works in markets, the most adversarial, most honest, fastest-feedback environment available to ordinary people, the broader architecture has a chance of being worth anything. The broadest claim comes last, not first.

The moat — ecosystem, context, and applied safety

The moat is process engineering. Not access to one model provider. Not access to one rail. Not a prettier interface. The moat is understanding how real work actually gets done in markets, the sequence from thesis to plan to execution to monitoring to review and encoding that understanding into systems trustworthy enough to operate with real capital. Process engineering is durable in a way that feature sets are not, because it requires domain knowledge that cannot be compressed into a foundation model and lived experience that cannot be abstracted into a product spec.

What compounds:

1. **The control plane is the product.** We are not trying to own every rail, every venue, or every primitive. We own the operator-facing orchestration layer that makes fragmented rails operable with real context, real approvals, and real trust. In an economy where infrastructure commoditizes, the layer that shapes operator behavior captures economics. *Owning the user does not mean owning distribution. It means owning the right to choose where the economics settle.* This is routing power: once multiple protocols, venues, and providers can satisfy the same operator job, the position that decides where to send the flow captures a share of every layer below it. Multi-venue breadth is not a feature list. It is a margin strategy.
2. **The seam is the moat.** DeFi will not replace TradFi. Neobanks will not replace brokers. Crypto exchanges will not replace banks. The real product is the integration layer where all of them coexist and where operators move value between contexts. Gordon is designed to own that seam.
3. **The data flywheel is personal, not generic.** Every operator builds a durable financial context: theses, trades, mistakes, outcomes, risk tolerance, preferred venues, discipline patterns. That context is the agent's working memory and the operator's portable asset. This is not training data for a general model. It is the operator's own operating state, a preference-economy asset in the specific sense named in the academic literature. In the preference economy, the highest-value position is the trusted aggregator that operators and other agents query before deciding. Gordon is designed to become that oracle for each operator: not a general model, but a personal financial judgment layer that compounds with every trade and grows harder to replicate with time.
4. **Ecosystem depth across Web3 and TradFi.** Gordon is built as an open, protocol-compatible workbench rather than a walled silo. That lets us integrate deeply into Web3 (Solana, EVM, Polkadot, Base, Chainlink, Uniswap, MoonPay, Helius) and TradFi (direct broker access plus the standard exchange APIs) at a depth no single vertical-only competitor can match. Technical success can coexist with economic disappointment, infrastructure players who built the rails that power entire ecosystems frequently watch their economics captured upstream by the operator-aligned layer. General Liquidity is designed to be that layer.
5. **Trading-native AI safety.** General Liquidity is an applied product and research lab, and a significant share of the research output is the safety infrastructure specific to autonomous trading. Red-teaming protocols that stress-test the agent against adversarial prompts, market manipulation attempts, and hostile market regimes. A **constitution for trading** that encodes the rules the agent must honor even when the operator is asleep, the market is broken, or the model is uncertain. That constitution is derived from the foundational literature of systematic risk management: Van Tharp's position sizing methodology, the Market Wizards consensus on the 1–2% per-trade risk rule (Eckhardt, Kovner, Seykota), fractional Kelly Criterion (0.25× for consumer use), Turtle Trader pyramiding discipline, and the correlation and concentration controls that Larry Hite and the institutional quant generation built over decades. Gordon's safety surface is not a product decision. It is the institutional trading floor's hard-won knowledge, encoded into a system that enforces it even when the operator forgets. Behavioral guardrails enforced across every venue call and every permission mode. A set of trading-specific AI safety principles we publish on the record. When the first autonomous-trading failure headlines hit, and they will, the category will bifurcate between operators who built safety from day one and everyone else. We are the first camp.
6. **Applied research cadence.** Our public writing is the research record, published in real time, that sets the vocabulary of the control-plane category before the product needs it.

Why Big Tech does not eat this

The obvious skeptic argument is *"won't OpenAI or Anthropic or Google or Amazon or Apple simply eat this wedge?"* The structure of the market gives a clean answer. The agentic economy is emerging in two tiers: a competitive application layer on top of a monopolistic infrastructure layer (foundation models, cloud, identity). The infrastructure layer consolidates around a handful of players who benefit *more* when the application layer above them is diverse, competitive, and vertical-deep. Every API call, every compute-hour, every token generated by vertical players routes revenue back into the infrastructure layer. Foundation-model labs do not want to run broker compliance, venue operations, or regulated financial integrations. That work is vertical, operational, relentless, and directly adjacent to our founder's decade of technical experience. The labs monetize our success. They do not compete with it.

The harder version of the skeptic argument is *"won't Apple / Google / Meta win through their device and identity stacks?"* This is McKinsey's direct concern. The platforms have pre-installed distribution, identity, payment, and increasingly conversational interfaces. Our answer is that the agentic economy is dividing into **agentic walled gardens**: Amazon's agent only shops Amazon, Apple Intelligence only recommends Apple Pay, Google's agent only routes to Google Travel and an open web of **operator-aligned agents** that serve the operator against the platforms. Gordon is deliberately non-aligned: no hidden revenue share with venues, no dark patterns, no OS lock-in. As the walled gardens accumulate structural conflicts of interest, demand for genuinely operator-aligned alternatives grows. This is the same logic that made Signal and ProtonMail viable against WhatsApp and Gmail, except the underlying product here is an order of magnitude more valuable per user.

Regulation is a tailwind, not a tax

We are starting in the **US first** and expanding globally. Every major financial-agent regulation emerging across the US, UK, and EU in 2025–2026 mandates exactly the architecture Gordon already ships.

- **US — CFPB Section 1033** (finalized 2024): standardizes consumer data portability and authorized third-party access. Enables agents to pull account data legitimately on the operator's behalf. Gordon is designed for exactly this access pattern.
- **US — NIST AI Risk Management Framework**: voluntary but increasingly load-bearing in enterprise procurement and in SEC and FINRA guidance on AI use in financial services. Gordon's mapping of model use, risk measurement, and policy management is native to how it is built.
- **US — state-level**: Colorado AI Act (2024), California's pending AI legislation, and expected federal executive-order-derived standards create a patchwork that any serious AI-in-finance product has to handle cleanly.
- **Global — EU AI Act** (full effect August 2026): a critical clarification here matters. Annex III of the EU AI Act does **not** classify algorithmic trading as high-risk: high-risk categories cover credit scoring, insurance underwriting, and biometric onboarding. Trading execution sits outside that classification. Where the Act does apply to Gordon, it requires explainability, auditability, permission scoping, and human override for any AI system touching financial decisions. Gordon's six-mode permission model, audit log, and constitution for trading are a pre-built compliance surface for any EU expansion and they position Gordon correctly on both sides of the Act: compliant where required, and explicitly unencumbered where the regulation leaves room.
- **Global — UK FCA** is moving in a similar direction, and Asia-Pacific regulators are watching both US and EU templates closely.

Competitors shipping under-instrumented AI chatbot layers on top of existing products will need to rebuild for these regulations. Gordon was built for them from day one.

Competitive landscape

We are not alone in this space, and that is the point. The pattern repeats: every structural shift in finance has an infrastructure side and a consumer side. The infrastructure side is already crowded with serious players. The consumer side is empty.

Player	Layer	Angle	Relationship to Gordon
Coinbase AgentKit + x402 + World ID	Venue-native infrastructure	Crypto rails for agents, identity for Sybil resistance. Public pivot to "Agentic Web."	Adjacent. We integrate with them. Coinbase's pivot validates the category.
Griffin (agentic-first banking core)	Banking infrastructure	BaaS for fintechs building agentic products	Adjacent. Potential substrate partner.
Human.tech Agentic WaaP	Custody and key management	Wallet-as-protocol with 2PC custody and policy engine. Identity stack protects \$500M+ from Sybil attacks, secured by \$3B+ restaked ETH.	Potential integration partner.
Visa / Mastercard / PayPal	Payment rails	Agent-initiated payments, mandates, orchestration	Rails we sit on top of.
NickAI (Galaxy Digital, launched March 12 2026)	Visual workflow builder for agentic trading	Galaxy-backed distribution; visual workflow approach targeting crypto-native traders	Different paradigm: visual workflow vs. chat-first operator terminal. No plan-first preview. No cross-venue coverage matching Gordon's depth. Galaxy backing means better distribution, not better architecture.
Public.com (launched "Agentic Brokerage," March 31 2026)	Retail brokerage + agentic layer	US equities only; execution delegation with no plan-first preview and no published permission model	Direct competitor in narrative, not in architecture. No crypto, no cross-venue, no structured diff before execution. Their failure mode is a trade the operator didn't see coming, Gordon's window.
3Commas / QuantPilot (QuantPilot launched April 2 2026)	Crypto bot platform + AI research layer	7+ year execution track record; QuantPilot adds AI strategy generation on top of existing DCA/grid/composite bots	Crypto-only. Bot architecture (fixed rules) vs. agent architecture (ongoing reasoning). History of API key exposure incidents. Research-to-execution coupling is loose: no plan preview, no permission modes.
Robo-advisors (Betterment, Wealthfront)	Consumer advisory	Algorithmic allocation, rebalancing	Strategy-first and chat-first is a different shape than static allocation.
DeFi copilots (aixbt, Olas, Wayfinder, Velvet)	Crypto-native agents	Onchain yield and trading	Narrow. Crypto-only. No stocks, no broker, no personal finance.
Neobanks with AI chat (Revolut, SoFi, Nubank)	Consumer banking	LLM chatbots on top of existing products	Chatbots are not agents. No real execution, no cross-venue, no trading surface.

TradFi brokers with AI (Robinhood, Schwab)	Retail brokerage	Generative AI for research summaries	Single-venue, single-market, no execution delegation.
Platform agents (Apple Intelligence, Google Assistant, Microsoft Copilot)	OS-level assistants	Aligned with platform interests	Conflict-of-interest risk creates demand for operator-aligned alternatives.

Nobody else is building the chat-first operator environment that spans crypto and stocks, carries real execution authority, ships with a real risk gate, is architected for the regulations that drop in 2026, and is designed from day one as the beginning of a personal finance operating system. Every competitor is either a crypto-only agent, a single-venue feature, an infrastructure play with no consumer surface, or a platform-aligned assistant with structural conflicts of interest. We sit at the seam.

--	--	--

The steelman — the strongest skeptic argument, answered

The strongest argument against us comes from Naeem Aslam ([Investing.com](https://www.investing.com)): "*Performance is not validated in days. It is revealed over cycles.*" AI trading systems that look good in stable conditions tend to fail during volatility, and an autonomous system executing real orders turns a flawed signal into a flawed action with direct financial consequences. Markets compress truth over years, not weeks. This is the serious version of the skepticism.

Three responses.

First, we agree with the premise and built the product around it. Gordon defaults to `ask` mode. `auto` mode exists but requires an explicit mandate and ships with risk gates, exposure caps, drawdown limits, and a kill switch. `paper` mode is a first-class citizen of the permission truth table, the safe on-ramp to live execution. We are not betting on agentic execution being validated in days. We are betting on a product that *lets the operator validate it over cycles* while keeping them in control of every cycle. That is the feature.

Second, the skeptic argument cuts deeper against the institutional side than against us. Institutional quant funds bet capital on models before out-of-sample cycles validate them. That is the whole point of trading. Our operators do not. Gordon's worst case in `ask` mode is that the operator rejects every trade and learns nothing. The institutional worst case is a fund blowing up. The system architecture shifts risk *away* from the operator, which is precisely the asymmetry retail investors currently suffer the other direction.

Third, the long-term case for agentic retail is stronger under Aslam's frame, not weaker. If performance is revealed over cycles, the product that wins is the one that keeps an operator in the game across cycles, preserving context, memory, discipline, and learning from mistakes. That is exactly the personal-context data flywheel Gordon is designed to build. Every competitor is optimized for today's trade. Gordon is optimized for an operator's decade of trades. Time is on our side.

10-year vision

In 2036, General Liquidity is the financial operating system for the agentic generation. Gordon, in its terminal form, its Desk form, and its personal-finance form, is how a material fraction of Gen Z and younger millennials actually operate their money.

- Tens of millions of operators have accumulated multi-year financial context inside Gordon. That context is the primary asset of the company: every plan, every trade, every mistake, every calibration made, portable and compounding.
- Gordon Desk is the default environment for any serious retail operator who wants to run strategies, backtests, and cross-venue portfolios. It is to personal trading what Cursor is to coding, the place where humans and agents collaborate at the level of an agentic Bloomberg Terminal, with shared context and real authority. Agentic multiplayer is a normal work pattern. Teams of two, three, or ten operators run shared theses together with shared agent infrastructure.
- Gordon agents transact, pay, negotiate, reconcile, and optimize across cards, banks, wallets, exchanges, payroll, and subscriptions. The boundary between trading and personal finance has dissolved because both are just forms of delegated economic action. Every external merchant, insurer, bank, and broker builds API-first, machine-legible products because Gordon, not a human, is increasingly the counterparty.
- The applied-research arm has set the vocabulary of the control-plane category in public, on the record, and has become a reference in how trading-native AI safety is defined and measured.
- The company is the canonical worldbuilding answer to the question *"what does it mean for ordinary people to have economic agency in the age of agents?"* The answer is visible, operable, profitable, and used.

If we are right about programmability without agency being the trap of the next decade, General Liquidity is the company that broke the trap for everyone who was not already on the right side of it.