



General Liquidity

Trade with discipline. Not willpower.

contact@generalliquidity.com

The thesis — worldbuilding

The economy is more programmable, more adversarial, and harsher at scale than at any point in modern history. Financial institutions already operate production-grade AI — quant systems, state-of-the-art economists, and microstructure models — positioned so whales can beat retail positions by design. AI has killed manual trading at the institutional end. Agents are eating the market. In parallel, the career ladder for young people is losing its bottom rungs. Cognitive entry-level work was the training ground for every professional career of the last century, and cognitive entry-level work is exactly what agents do best. What analysts call the hollowing out of the corporate structure is the same phenomenon this generation is living through as *"there are no jobs and the ones that exist don't lead anywhere."*

There is a wealth transfer happening right now, caused by AI itself, and the only question is whether ordinary people get to participate or get to watch. Research from Allianz already shows that software investment produces the strongest labor-substitution effect of any capital category — a **0.22–0.29% labor reduction per 1% increase in software investment**, with finance and real estate facing the steepest employment declines. AI agents are a new factor of production, neither classical labor nor classical capital, what the systemic analyses are starting to call **agentic capital** or **digital labor**. The returns to that factor flow to whoever owns it. Right now that is the platform layer and the institutional trading floor. It does not have to stay that way.

The peer-reviewed literature is catching up to what the institutional floor already knows. The Google DeepMind *Virtual Agent Economies* paper (arxiv 2509.10147, 2025) names the mechanism precisely: *"More capable AI assistants negotiate better deals for their users."* The implication is not subtle. In every market where AI agents transact on their principal's behalf, the quality of the agent determines the deal. A retail

investor with a 2016-era chatbot recommendation is not competing with a hedge fund's full-stack execution infrastructure. They are being extracted from. The paper further documents the systemic risks: flash-crash dynamics from correlated agentic behavior, and the critical finding that multi-tiered oversight frameworks — exactly the kind Gordon's six-mode permission model implements — are the primary known mitigation for agentic market instability. The academic literature and Gordon's architecture are pointing at the same design.

The participation gap is not about access. Most people who will benefit from Gordon already have a brokerage account and a crypto wallet. They are inside the system. The problem is that they are outside the loop that matters: the loop where a thesis becomes a plan, a plan becomes a position, a position is monitored, and outcomes compound into judgment. Having an account is not the same as operating an account. Participation is not agency.

General Liquidity exists to make sure the agentic generation gets to participate, not just watch. We are building **agentic vibe trading**: a consumer AI product category in which a person with no quant background, no algo experience, no economist pedigree, and no finance degree can operate in markets with the same control plane institutions already use. Trading is the wedge because it is where wealth gets extracted from markets today. After that we extend into personal finance, and the context each operator accumulates (theses, trades, preferences, discipline, outcomes) compounds into a **personal financial mobility engine**.

Programmability without agency is a bigger trap than no programmability at all. More rails, more models, more access, and a population that understands less and less of what is happening to their money. The field's most common responses to this problem all collapse toward the wrong answer. Foundation models absorbing the application layer — they commoditize the model, they do not own the operator context built on top of it. New rails making legacy infrastructure obsolete — DeFi and TradFi will coexist for decades; the productive position is the seam between them, not a bet on migration. Natural language solving the software problem — natural language is an interface, not a control plane; a chatbot that reasons aloud but cannot execute is a more persuasive version of the same access problem, not an answer to it. In each case the surface changes. The asymmetry remains.

The control plane between raw capability and operator intent is the product. **Access is not agency. Raw rails do not create trust. Raw models do not create judgment.** That plane is what General Liquidity builds.

The opportunity — the agentic economy is happening now

AI is projected to contribute \$15 trillion to the global economy by 2030. Agentic AI specifically is projected to add \$2.6–4.4 trillion globally over the next 2–5 years (Allianz). Those are macro numbers. The financial services layer has its own.

Global payments industry annual revenue

\$2.7T

McKinsey

Consumer credit-card revenue (2024)	\$234B	McKinsey
Global retail-bank net interest income share of profit	~30%	McKinsey
Addressable autopilot-services market (aggregate)	\$1.5T+	Sequoia
Services budget per dollar of software budget	6:1	Sequoia
Agentic AI market (2025 → 2033)	\$4.54B → \$98.26B	Grand View Research

That last line deserves a number of its own: **46.87% CAGR over 8 years** is not a category trend. It is a generational platform shift. The CFA Institute confirms that 73% of Y Combinator investments between January 2024 and June 2025 were in agentic AI companies. Earnings call mentions of AI agents increased 4× quarter-over-quarter in Q4 2024 alone. The capital is already flowing; the infrastructure buildout is happening in real time.

Sequoia's autopilot opportunity map breaks the services number into specific verticals, each its own trillion-dollar-class pattern.

Insurance brokerage	\$140–200B	WithCoverage, Harper
Claims adjusting	\$50–80B	Pace, Strala
Healthcare revenue-cycle	\$50–80B	Anterior
Accounting & audit	\$50–80B	Rillet, Basis
Tax advisory	\$30–35B	TaxGPT, Skalar, Ravical
Legal transactional	\$20–25B	Harvey, Crosby, Lawhive
IT managed services	\$100B+	Edra, Serval
Supply chain / procurement	\$200B+	Magentic, AskLio, Tacto
Recruitment & staffing	\$200B+	Juicebox, Mercor, Jack & Jill
Management consulting	\$300–400B	early — TBD

Gordon's wedge — delegated trading for retail operators — is the consumer-finance analog of every line on that table. Intelligence-heavy, currently outsourced (to brokers, advisors, influencers, or sheer luck), and waiting for the first credible autopilot to collapse the cost of doing it well.

Adoption is no longer theoretical on either side. Institutional, retail, and regulatory signals are converging inside a 12-month window:

- **Gartner:** 33% of enterprise software will include agentic AI by 2028, up from under 1% in 2024. 15% of day-to-day work decisions autonomous by 2028, from 0%.
- **PagerDuty:** 51% of companies already have agents in production. 78% plan adoption.
- **Goldman Sachs:** 300M FTE jobs globally exposed to generative-AI automation.
- **Yu'e Bao** (China, money-market fund app) moved \$268B of consumer cash in under four years once the rate gap became visible. Consumers do move money when they can see the opportunity.
- **Raisin** (Europe, deposit marketplace) has \$80B across 250 banks. Same thesis, different geography.
- **Santander + Mastercard** completed Europe's first live end-to-end AI-executed regulated financial transaction in March 2026. Not a pilot. A live payment with delegated authority.
- **DTCC** reports its Risk Calculator AI Assistant runs at ~97% accuracy, cutting client-meeting prep from one week to one day and driving a fourfold productivity lift for client-facing teams. This is the plumbing of global capital markets publicly admitting agents already work at scale.
- **Moody's** reports its AI Research Assistant users consume 60% more research while cutting task completion times by 30%.
- **Coinbase** is publicly pivoting to become mission-critical infrastructure for the Agentic Web — AgentKit, x402 with Cloudflare, World ID for Sybil resistance, Base + USDC as agent settlement rails. Crypto increasingly functions as machine-native financial infrastructure — programmable, composable, and built in a form software can use natively.
- **Human.tech's Agentic WaaP** (Wallet-as-a-Protocol) launched at WalletCon 2026. Its identity infrastructure protects \$500M+ from Sybil attacks, secured by \$3B+ in restaked ETH.
- **MIT + Stanford** field study on AI augmentation: 14% productivity lift on average, 34% for less-experienced workers. AI is a great equalizer. The biggest lift goes to the people furthest from the frontier. That is the empirical foundation of the economic-mobility thesis.
- The operational plumbing is ready. FedNow (US), TIPS and SEPA Instant (EU), and UPI (India) have eliminated settlement lag. Stablecoin rails (Visa-USDC, Circle, Coinbase Payments) add programmable, sub-second transfers. The rails an agent needs to sweep funds, split orders, or pay a vendor in real time are live today.

The consumer side is still mostly empty. The institutional rails are shipping; the retail control plane is not. That is the gap.

The paradigm shift we ride

Four authoritative voices have published the same thesis inside the last 18 months, each pointing at the exact wedge General Liquidity is building. When that many credible voices converge on the same claim in the same year, it is a market forming.

Sequoia — Services, the New Software (Julien Bek, March 2026). *"If you sell the tool, you're in a race against the model. But if you sell the work, every improvement in the model makes your service faster, cheaper, and harder to compete with."* Copilots to autopilots. Gen Z and younger millennials already outsource financial

judgment: to influencers, to TikTok, to Robinhood's default UX, to sheer guesswork. Replacing that informal outsourcing with Gordon is a clean vendor swap. *"Today, more tasks are started by agents than by humans."* The same inflection is arriving in retail finance. *"The work budget in any profession dwarfs the tool budget, and autopilots capture the work budget from day one."* Sequoia also draws the line between intelligence and judgment: executing trades is intelligence work, and agents do it better. Knowing which thesis to deploy, in what market regime — that is judgment work, and it still belongs to the operator. Gordon's architecture is this split made concrete.

Block — *From Hierarchy to Intelligence* (Jack Dorsey and Roelof Botha, March 2026). *"A world model that can't touch the world is just a database."* The four-layer model — capabilities, world model, intelligence layer, interfaces — is the pattern Gordon is built around. Block also makes the most direct argument for why financial behavior data is uniquely valuable: *"Money is the most honest signal in the world. People lie on surveys. They ignore ads. They abandon carts. But when they spend, save, send, borrow, or repay, that's the truth."* Gordon's per-operator financial context is built on this signal. Not aggregate behavioral proxies. The operator's own revealed preferences, compounding under the pressure of real capital over real time.

Pear VC — *Request for Startups* (S26 cohort). Three of Pear's twelve RFS categories are directly Gordon-shaped. **Category 1, *Plan mode for every knowledge worker*** — Gordon has a literal `plan` permission mode and a plan-first operator workflow that brings structured planning to financial decision-making. Pear's framing: current AI tools produce outputs that are "60–70% of the way there" — enough to generate the analysis, not enough to close the trade. **Category 2, *Financial-grade agent infrastructure*** — trust, policy, and execution layers for agents operating in regulated finance. **Category 8, *Intent-native commerce*** — agents as first-class buyers, programmatic pricing, catalogs as APIs.

McKinsey — *The End of Inertia* (August 2025). *"Margin accrues in the gap between what consumers could do and what they actually do. Agentic AI breaks that equation."* McKinsey names five control points in the agentic economy: credentialing and identity, trust and liability wrappers, merchant and platform integration, decisioning logic, and behavioral data with intent signals. Gordon owns the operator-facing surface of all five. *"Relevance must be re-earned moment by moment — every time an agent refreshes its rankings."* Platform-aligned agents that depend on brand inertia will fail this test. Gordon is designed to pass it by construction.

The Agentic Economy (arxiv 2505.15799). The paper argues the central economic shift is from the attention economy to the **preference economy**. Gordon does not sell attention. It sells operator preference data that compounds into a durable personal financial context. The flywheel works on subscriptions alone, and the data moat is a fundamentally different asset class from the one consumer social products built.

Google DeepMind — *Virtual Agent Economies* (arxiv 2509.10147). The empirical finding that more capable agents negotiate better outcomes for their principals has a direct policy implication: agent capability gaps between retail and institutional principals are not self-correcting. Left unchecked, they widen. The paper further argues that multi-tiered oversight frameworks are the primary known mitigation for correlated agentic behavior creating flash-crash dynamics. Gordon's six-mode permission model is a concrete implementation of that framework — not as a regulatory checkbox, but as the product feature that makes real delegation safe.

The operational vocabulary converging across these sources is **human above the loop**, not *human in the loop*

(Pawel Gmyrek, ILO, via WEF 2024). *In the loop* implies the human is a step in the pipeline. *Above the loop* implies the human sets the objective and supervises. Gordon's six-mode permission model is the cleanest implementation of human-above-the-loop currently shipping in consumer finance.

The problem — the agency bottleneck

Retail finance has been sold a story that access equals empowerment. Commission-free brokerages, neobanks, DeFi wallets, crypto exchanges — all expanded access, none created agency. A person with a Robinhood account and a Coinbase account participates in markets without operating them. They can click; they cannot reason. When prices move, they react. The tools tell them *what* is happening and leave the *what now* entirely on the user. A user who usually has no finance background, no economics training, no access to institutional research, and no infrastructure to turn a market view into a disciplined position.

Institutions, meanwhile, run fully systematized operations: research, thesis, sizing, execution, risk, review, compounding. Renaissance Technologies, Jane Street, HRT, Citadel — the top tier has iterated on systematic execution for decades. An institutional trader has agents managing positions while a 24-year-old with a \$40k salary is squinting at a candlestick chart on their phone. *The asymmetry is the product financial institutions sell.* SP Global's analysis of market fragmentation confirms that the proliferation of venues has not helped retail traders navigate — it has helped institutional arbitrageurs extract from the fragmentation.

The numbers make it precise. McKinsey documents the specific cost of inertia: US banks paid 0.07% on checking accounts and 0.38% on savings accounts in 2024, while online savings rates exceeded 4%.

Consumers absorb more than three percentage points of lost yield every year simply because switching is friction-heavy enough to be avoided. The execution gap in trading is worse. 60–75% of US equity volume is algorithmic. Only 7–13% of human Polymarket traders achieve positive P&L, versus 37%+ for AI agents on the same markets. The gap is not closing. It is widening with every improvement in institutional AI.

The Yu'e Bao and Raisin evidence shows what happens when friction is removed — \$268B and \$80B moved, respectively, once the gap became visible and acting on it required no effort. Gordon's *ask* mode and plan-first execution is the architectural answer: the gap between what the operator should do and what they actually do is closed not by financial education but by reducing the cost of acting on what they already know.

What is changing is that the tools necessary to close this asymmetry finally exist, and they create a governance problem the industry is only starting to name. Hanah-Marie Darley (Geordie AI): *"Permissions and governance must travel with the agent — not sit in static policy documents. The rules that define what the agent is allowed to do need to travel with it as it moves across tools, services, and markets."* Every retail fintech shipping AI today is failing this test. LLM chatbots that *recommend* without executing. Bolt-on credentials that do not carry policy context. Gordon's six-mode permission truth table, enforced at every stage of a decision, is an architectural answer to Darley's frame, built before the milestone forced the industry to articulate it.

The academic literature adds a second taxonomy that clarifies the competitive landscape. Most current AI

products are *siloed* service agents: Amazon's Rufus operates only inside Amazon, Expedia's Romie operates only inside Expedia. They cannot carry context across venue boundaries or represent the operator's full picture. Gordon is an *end-to-end* agent: cross-venue, cross-market, persistent context that travels with the operator. The literature further distinguishes *unscripted* interactions (technically possible via natural language) from *unrestricted* interactions (sanctioned by governance and market structure). Natural language gets to unscripted. The permission model, the audit log, and the risk gates get to unrestricted. Most retail AI products are unscripted. Gordon is designed to be both.

Agent commerce — and by extension, agent-delegated finance — splits the world into tourists who treat every transaction as an isolated one-off, and locals who operate inside supplier relationships, negotiated terms, and accumulated context. Retail investors are tourists in markets today. Gordon turns them into locals.

Agentic AI is the first technology that can close the institutional-retail asymmetry, and it is closing it on the wrong side first. OpenAI Operator, Manus AI, Perplexity, and Coinbase AgentKit are already giving sophisticated users agentic tooling. The institutional end is accelerating faster than the retail end. If nothing intervenes, the wealth transfer AI is causing will skew exactly the way every prior technological wealth transfer has skewed: toward capital, away from labor, away from the young.

The customer — who Gordon is built for

Gordon is a direct consumer product. The person who pays is the person who uses it. Understanding who that person is — and who it is not — determines every product, pricing, and GTM decision.

The Stuck Operator (primary ICP)

THE STUCK OPERATOR
Age: 25-32
Occupation: Software engineer / tech worker / finance adjacent
Income: \$60-100k
Portfolio: \$15-40k actively managed
Location: US (SF, NYC, Austin, remote)
Tech comfort: High - uses Claude Code, Cursor, or CLIs
DAILY REALITY:
• Reads Cobie, Avi Bhatt, Murad, or macro Twitter in the AM
• Has a thesis. Usually a good one.
• Opens Robinhood or Coinbase, hesitates, maybe trades wrong
• Has tried: ChatGPT + Alpaca API, Pine Script, 3Commas
• All of them either broke, were too complex, or didn't work the way they meant
PAINS (ranked by score - urgency × cost × frequency):
1. Emotional execution failure - "I know what I should do.

I can't make myself do it." (score: 11/12)
2. Execution-systematization gap – "Has thesis, can't build infrastructure." (score: 11/12)
3. Tool fragmentation – "I use 4 apps and none talk to each other." (score: 10/12)
4. Copy trading failure – "Gave up control. Still lost."
5. Failed DIY automation – "Built something. It broke."
DESIRED OUTCOME:
"Something that enforces my own rules on me, so I trade like I know how to trade, not like I feel in the moment."
SWITCHING TRIGGERS (what makes them look now):
• Recent revenge trade or FOMO entry they saw coming
• DIY automation attempt that broke
• Copy trading burn – lost money with no sense of control
• Watched a thesis play out but couldn't capture it cleanly

What they need to believe before paying:

1. "This actually executes — it's not just another chatbot."
2. "It won't do something I didn't ask for."
3. "I can start on paper and move to live when I'm ready."
4. "It covers both my crypto and my stocks."

The primary pain to lead with: *"You already know what to do. Gordon makes sure you actually do it."*

Discipline is the product. Not alpha.

The Crypto-Native Power User (secondary — first 100 operators)

Age 22–30, Web3 developer or DeFi researcher. Uses Binance Advanced, Hyperliquid, and Uniswap simultaneously. Monitors multiple chains and DEX flows. Has custom alerts and possibly a home-built bot. Has missed entries because monitoring is fragmented.

Pain: *"I monitor 6 things to trade one thing."* and *"There's no product that does onchain AND CEX AND stocks."*

Why they convert first: Lowest time-to-value. Already has accounts and API credentials. Already feels the fragmentation pain acutely. Most likely to become champions.

Sophisticated Late Majority (Gordon Desk era, ~18 months)

Age 30–45, finance professional or entrepreneur, \$100k–\$500k managed. Needs visual interface (Gordon Desk), institutional-grade audit log. Will pay \$200–\$300/month once the category is proven. Not the target now — CLI interface and friends-alpha signal quality are wrong for their risk tolerance.

Anti-personas (who NOT to target in Phase 1)

Passive index investors	No active trading intent; feature set is overkill
Complete beginners	Time-to-value too long; product amplifies knowledge, doesn't generate it
Professional fund managers	Regulatory constraints, liability — requires institutional-grade product
Pure DeFi/NFT speculators	Want fully autonomous agents; will fight the safety model
"Get rich quick" traders	High churn, high support cost, misaligned with thesis-driven discipline

Where to find them

- r/algotrading, r/CryptoCurrency, r/stocks, r/options, r/wallstreetbets (power users)
- Finance Twitter/X (Avi, Murad, Cobie followers)
- Crypto Discord servers (DeFi Research, Bankless community, Messari)
- Hacker News (Show HN for the CLI launch)
- GitHub (developers who star trading bot repos)
- Claude Code / Cursor community (highest-fit technical profile)

The product — Gordon, and the product ladder

Gordon is the first agentic vibe trading product. It operates across crypto exchanges and stock brokers, runs a chat-first operator workflow, and carries real execution authority under a human-controlled permission model.

- **Chat-first operator workflow.** Scan, analyze, plan, preview, approve, execute, monitor, reconcile. No dashboards, no cockpit, no twelve tabs.
- **Multi-venue by design.** 9 crypto exchanges (Binance, Coinbase, Kraken, Bitfinex, Gemini, OKX, Hyperliquid + more) and 9 stock brokers (Alpaca, Schwab, Interactive Brokers, E*TRADE, Tastytrade, TradeStation, Tradier, Trading 212, Webull), plus onchain execution across Solana, EVM, Polkadot, and Base.
- **Deep research and analysis.** Market discovery, multi-timeframe technical analysis, onchain and DEX flows, fundamentals across stocks, ETFs, indices, bonds, crypto, and macro, SEC filings, news entity extraction, plus a backtesting engine with Monte Carlo and walk-forward validation.
- **Six permission modes.** auto, ask, strict, paper, observe, plan. A single truth table that lets the operator dial autonomy from read-only research all the way to fully autonomous execution within risk gates. The same policy is enforced before, during, and at every venue call. This is the "permissions travel with the agent" pattern in shipping form.
- **Plan-first execution.** Every trade is previewed as a structured diff before anything hits a venue. Bracket, grid, DCA, and OCO plans with trailing stops and partial take-profits. No order leaves without a clean

reason and an explicit approval step under the operator's mode of choice.

- **Real safety surface.** Input guardrails, access control, rate limits, risk gates, and full audit logging. A system an operator can actually trust with capital.

The product ladder

1. **Gordon (now).** Agentic vibe trading for serious operators. The Claude Code / Codex of vibe trading: terminal-native, keyboard-sovereign, scriptable, auditable. The proving ground.
2. **Gordon Desk (next, ~12–18 months).** The visual workspace that brings research, planning, execution, monitoring, and review into one environment. Think *agentic Bloomberg Terminal × Cursor × Claude Projects* for markets. Multi-agent, multi-operator, shared context. **Agentic multiplayer** — where shared theses, collaborative backtesting, and team trading become normal.
3. **Gordon Personal Finance OS (later, 24–36 months).** The post-wallet, post-broker financial home. Gordon transacts on the operator's behalf across cards, banks, wallets, exchanges, payroll, tax, and subscriptions. Every operator's financial context — theses, trades, preferences, discipline, outcomes — becomes a compounding personal asset that the agent operates from. *"The wallet opened the door. It did not finish the room."*

Gordon CLI is not a waypoint on the way to Gordon Desk. It is the test. If delegation works in markets — the most adversarial, most honest, fastest-feedback environment available to ordinary people — the broader architecture has a chance of being worth anything. The broadest claim comes last, not first.

Market positioning

The onliness statement

Gordon is the only agentic trading terminal that executes across crypto and equities with a plan-first permission model — so the operator always knows exactly what it will do before it does it.

Six-element form:

- **WHAT:** Agentic trading terminal for retail operators
- **HOW:** Plan-first structured diff preview + six-mode human-controlled permission model enforced at every venue call
- **WHO:** Self-directed operators with market conviction but no execution infrastructure (The Stuck Operator)
- **WHERE:** Crypto exchanges + stock brokers + onchain, simultaneously
- **WHY:** To convert market knowledge into disciplined, monitored positions without emotional interference
- **WHEN:** As institutional AI widens the execution gap and "vibe trading" enters mainstream vocabulary

This onliness holds on three simultaneous axes: multi-venue TradFi+DeFi depth, plan-first preview, and six-

mode permission model. No single competitor matches all three.

The positioning statement (Moore)

For self-directed retail operators who have market conviction but no infrastructure to execute it consistently, **Gordon is an** agentic trading terminal **that** converts your thesis into a managed, monitored position across crypto and stocks — showing you exactly what it will do before it does it. **Unlike** single-venue AI features, copy-trading platforms, or passive robo-advisors, **Gordon** gives you the same control plane institutions already use: plan-first execution, six permission modes, and a research-to-execution workflow that enforces your discipline for you.

Tagline

Primary: *Trade with discipline. Not willpower.* — captures the core value theme (discipline infrastructure), speaks the customer language directly, works across audiences.

Secondary / subline: *The only trading terminal that shows you the plan before executing it.* — functional differentiator for the product page.

Community frame: *Pioneering agentic vibe trading.* — "vibe trading" is named and gaining vocabulary traction. Forbes: *"will be the 2027 standard for all brokerages."* The frame is named but unclaimed by any funded player. Every piece of Gordon content should reinforce this vocabulary before a larger competitor picks it up.

Competitive alternatives map

Public.com AI Agents (launched Mar 31, 2026)	Direct	Execute intent-driven trades, US equities	US brokerage only, no crypto exchanges, no plan-first preview, no backtesting
NickAI (Galaxy Digital-backed, Mar 2026)	Direct	Visual workflow trading across crypto + equities	Workflow-builder UX, no plan preview, developer-oriented
3Commas QuantPilot (Apr 2026)	Direct	Add research/strategy layer to bot trading	Crypto-only, legacy bot performance issues, no equities
Robinhood Cortex	Adjacent	AI-powered research and market analysis	Research summaries only, no execution delegation, single venue
TradingView + manual execution	Adjacent	Research, chart analysis, technical setup	No agent layer, purely manual execution, no cross-venue

Discord alpha groups + copy trading	Adjacent/ Manual	Capture signals, follow smart money	Structurally broken P&L, no risk gates, emotional dependency
Robo-advisors (Betterment, Wealthfront)	Adjacent	Automated passive allocation	Static, no thesis-driven execution, wrong mental model
Status quo (doing nothing)	Primary competitor	Stay informed	No execution, no systematization, pure reaction

The real competitor is the status quo. Most Stuck Operators do nothing systematic at all. Gordon competes with inaction as much as any named product.

Where Gordon wins vs. each alternative

vs. [Public.com](#): Multi-venue breadth (crypto + equities vs. US equities only), plan-first preview (Gordon only), permission model architecture (Gordon only), onchain execution (Gordon only). [Public.com](#) has brand momentum and institutional backing; Gordon has product depth.

vs. 3Commas: Agent vs. bot mental model — Gordon reasons about market conditions, 3Commas executes preset conditions. Gordon covers equities (3Commas does not). Gordon's trust architecture addresses the API key security issues that have historically plagued 3Commas.

vs. TradingView: Gordon is not a replacement for TradingView — it is the execution layer downstream of it. Operators who use TradingView for research can use Gordon to execute on that research without switching to a different product or manually replicating their analysis.

vs. copy trading: Gordon puts the operator above the loop. Copy trading puts the operator below it. The key insight: *copy trading failed because it gave up control AND still lost money*. Gordon preserves control.

vs. status quo: Paper mode eliminates the primary adoption barrier (fear of autonomous trades with real capital). The free tier delivers full research and paper trading value with zero capital risk — not a crippled demo, but a complete on-ramp.

The moat — ecosystem, context, and applied safety

The moat is process engineering. Not access to one model provider. Not access to one rail. Not a prettier interface. The moat is understanding how real work actually gets done in markets — the sequence from thesis to plan to execution to monitoring to review — and encoding that understanding into systems trustworthy enough to operate with real capital. Process engineering is durable in a way that feature sets are not, because it requires domain knowledge that cannot be compressed into a foundation model and lived experience that cannot be abstracted into a product spec.

What compounds:

1. **The control plane is the product.** We are not trying to own every rail, every venue, or every primitive. We own the operator-facing orchestration layer that makes fragmented rails operable with real context, real approvals, and real trust. In an economy where infrastructure commoditizes, the layer that shapes operator behavior captures economics. *Owning the user does not mean owning distribution. It means owning the right to choose where the economics settle.* This is routing power: once multiple protocols, venues, and providers can satisfy the same operator job, the position that decides where to send the flow captures a share of every layer below it. Multi-venue breadth is not a feature list. It is a margin strategy.
2. **The seam is the moat.** DeFi will not replace TradFi. Neobanks will not replace brokers. The real product is the integration layer where all of them coexist and where operators move value between contexts. Gordon is designed to own that seam. No competitor spans 9 crypto exchanges + onchain (Solana, EVM, Polkadot, Base) + 9 stock brokers simultaneously. This is a multi-year integration investment. Competitors have ≤ 5 venues. Each new venue connection creates incremental switching cost once operators are connected.
3. **The data flywheel is personal, not generic.** Every operator builds a durable financial context: theses, trades, mistakes, outcomes, risk tolerance, preferred venues, discipline patterns. That context is the agent's working memory and the operator's portable asset. This is not training data for a general model. It is the operator's own operating state — a preference-economy asset in the specific sense named in the academic literature. Per-operator context compounds: an operator using Gordon for 6 months has a persistent record that makes Gordon dramatically more useful to them than any new entrant. The moat is time-gated, which means it gets stronger, not weaker, as the product ages. A competitor launching today cannot retroactively acquire that operator history.
4. **Ecosystem depth across Web3 and TradFi.** Gordon is built as an open, protocol-compatible workbench rather than a walled silo. That lets us integrate deeply into Web3 (Solana, EVM, Polkadot, Base, Chainlink, Uniswap, MoonPay, Helius) and TradFi (direct broker access plus standard exchange APIs) at a depth no single vertical-only competitor can match. Technical success can coexist with economic disappointment — infrastructure players who built the rails that power entire ecosystems frequently watch their economics captured upstream by the operator-aligned layer. General Liquidity is designed to be that layer.
5. **Trading-native AI safety.** A significant share of our research output is the safety infrastructure specific to autonomous trading. Red-teaming protocols that stress-test the agent against adversarial prompts, market manipulation attempts, and hostile market regimes. A **constitution for trading** that encodes the rules the agent must honor even when the operator is asleep, the market is broken, or the model is uncertain. When the first autonomous-trading failure headlines hit — and they will — the category will bifurcate between operators who built safety from day one and everyone else. We are the first camp.
6. **Applied research cadence.** Our public writing is the research record, published in real time, that sets the vocabulary of the control-plane category before the product needs it.

The Can't/Won't test

Can't — The data flywheel has a time tax. Per-operator context cannot be retroactively acquired. A competitor launching tomorrow starts from zero operator history. The moat is being built today.

Can't — The seam between DeFi and TradFi is genuinely hard. 9 crypto exchanges + onchain + 9 stock brokers in one execution environment is a multi-year investment in venue-specific normalization, credential management, and permission enforcement across heterogeneous APIs. It is not a weekend integration.

Won't — Incumbents are structurally prevented from copying the permission model. Robinhood, Coinbase, and Schwab cannot build a plan-first preview that adds friction before execution — because they need the trade to execute to generate PFOF revenue. Their business model requires frictionless execution. Gordon's subscription model means adding deliberate friction is economically fine.

Won't — "Human-above-the-loop" conflicts with fully autonomous agent theses. A well-funded competitor building fully autonomous trading agents will not copy Gordon's permission model — it contradicts their product thesis. The market will bifurcate. Gordon should lean into the distinction, not blur it.

Why Big Tech does not eat this

The agentic economy is emerging in two tiers: a competitive application layer on top of a monopolistic infrastructure layer (foundation models, cloud, identity). The infrastructure layer benefits *more* when the application layer above it is diverse, competitive, and vertical-deep. Foundation-model labs do not want to run broker compliance, venue operations, or regulated financial integrations. That work is vertical, operational, and relentless. The harder version — *"won't Apple / Google / Meta win through their device and identity stacks?"* — has a structural answer: the agentic economy is dividing into **agentic walled gardens** (Amazon's agent only shops Amazon, Apple Intelligence only recommends Apple Pay) and an open web of **operator-aligned agents** that serve the operator against the platforms. Gordon is deliberately non-aligned: no hidden revenue share with venues, no dark patterns, no OS lock-in. As the walled gardens accumulate structural conflicts of interest, demand for genuinely operator-aligned alternatives grows. This is the same logic that made Signal and ProtonMail viable against WhatsApp and Gmail — except the underlying product here is an order of magnitude more valuable per user.

Regulation is a tailwind, not a tax

Every major financial-agent regulation emerging across the US, UK, and EU in 2025–2026 mandates exactly the architecture Gordon already ships.

- **US — CFPB Section 1033** (finalized 2024): standardizes consumer data portability and authorized third-party access. Enables agents to pull account data legitimately on the operator's behalf.
- **US — NIST AI Risk Management Framework:** voluntary but increasingly load-bearing in SEC and FINRA guidance on AI use in financial services. Gordon's mapping of model use, risk measurement, and policy management is native to how it is built.
- **Global — EU AI Act** (full effect August 2026): classifies financial-agent systems as high-risk. Requires explainability, auditability, permission scoping, and human override. Gordon's six-mode permission model and audit log are a pre-built compliance surface for any EU expansion. Competitors shipping

under-instrumented AI chatbot layers will need to rebuild for these regulations. Gordon was built for them from day one.

- **Global — UK FCA** is moving in a similar direction, and Asia-Pacific regulators are watching both US and EU templates closely.

Competitive landscape

We are not alone in this space, and that is the point. The pattern repeats: every structural shift in finance has an infrastructure side and a consumer side. The infrastructure side is already crowded with serious players. The consumer side is empty.

Coinbase AgentKit + x402 + World ID	Venue-native infrastructure	Crypto rails for agents, identity for Sybil resistance	Adjacent. We integrate with them. Their pivot validates the category.
Griffin (agentic-first banking core)	Banking infrastructure	BaaS for fintechs building agentic products	Adjacent. Potential substrate partner.
Human.tech Agentic WaaP	Custody and key management	Wallet-as-protocol with 2PC custody and policy engine	Potential integration partner.
Visa / Mastercard / PayPal	Payment rails	Agent-initiated payments, mandates, orchestration	Rails we sit on top of.
Robo-advisors (Betterment, Wealthfront)	Consumer advisory	Algorithmic allocation, rebalancing	Wrong mental model — static allocation vs. thesis-driven.
DeFi copilots (aixbt, Olas, Wayfinder, Velvet)	Crypto-native agents	Onchain yield and trading	Narrow. Crypto-only. No stocks, no broker.
Neobanks with AI chat (Revolut, SoFi, Nubank)	Consumer banking	LLM chatbots on top of existing products	Chatbots are not agents. No real execution, no cross-venue.
TradFi brokers with AI (Robinhood Cortex, Public, Schwab)	Retail brokerage	Generative AI for research summaries	Single-venue, single-market, no execution delegation.
Platform agents (Apple Intelligence, Google Assistant)	OS-level assistants	Aligned with platform interests	Conflict-of-interest risk creates demand for operator-aligned alternatives.

Feature matrix (selected)

Public.com					Cortex
Multi-venue crypto + stocks	Strong	Weak (US only)	Weak (crypto only)	Missing	Missing
Plan-first preview	Strong	Missing	Missing	Missing	Missing
Six-mode permission model	Strong	Missing	Missing	Missing	Missing
Chat-first workflow	Strong	Adequate	Weak	Missing	Adequate
Backtesting engine	Strong	Missing	Adequate	Adequate	Missing
Onchain execution	Strong	Missing	Missing	Missing	Missing
Paper / simulation mode	Strong	Unknown	Adequate	Adequate	Missing
Cross-venue operator context	Strong	Missing	Missing	Missing	Missing
Audit logging	Strong	Unknown	Weak	Missing	Missing

Nobody else is building the chat-first operator environment that spans crypto and stocks, carries real execution authority, ships with a real risk gate, is architected for the regulations that drop in 2026, and is designed from day one as the beginning of a personal finance operating system.

Business model & unit economics

Revenue tiers

Tier	Price	Access	Purpose
Free	\$0/month	Observe + paper modes, full research + backtesting, no live execution	Trust engine and acquisition funnel. Operator builds context during paper mode — free tier is complete for its scope, not a crippled demo.
Pro	\$49/month	Live execution, 3 venues	Primary Phase 1 revenue driver. Priced at "less than one emotional trading mistake per month."

Power	\$149/ month	Auto mode, all venues, full backtesting pipeline	Advanced operators who have validated strategies. 3× Pro ARPU.
Desk	\$299/ seat/ month	Multi-operator, shared context, agentic multiplayer	Gordon Desk era (~18 months). Enterprise-adjacent. Highest LTV per account.

Pricing anchors: TradingView Pro (\$59.95/month, no execution), 3Commas Pro (\$59/month, crypto-only bots), Cursor Pro (\$20/month) — establish the price band. Gordon's Power tier (\$149) consolidates and executes what this stack does piecemeal for ~\$120/month while delivering cross-venue execution Gordon's competitors cannot.

The free tier is not a compromise. It is the trust engine. An operator who runs paper mode for 4 weeks has built a working model of how Gordon behaves with their capital. When they upgrade to live execution, they are upgrading from a known system.

Unit economics

Metric	Pro Tier	Power	Notes
ARPU	\$49/month	\$149/ month	Blended target: \$65–85 at steady-state tier mix
Gross Margin	~91%	~92%	Primary COGS: LLM API (\$2-4/month Pro, \$6-12/ month Power)
Monthly Churn (target)	≤ 4.5%	≤ 2.5%	Discipline infrastructure creates stickiness
LTV (base case)	~\$912	~\$5,066	$LTV = ARPU \times (1/churn) \times \text{gross margin}$
CAC (organic / PLG)	~\$20–40	~\$20–40	GitHub + HN + Finance Twitter; near-zero marginal cost
CAC (blended, Phase 1)	~\$60–80	~\$60–80	Assumes community-dominant GTM
LTV:CAC (base case)	~11–15×	~63–84×	Extremely strong if PLG activation holds
CAC Payback	~1–2 months	< 1 month	Subscription product pays back fast at low CAC
Burn Multiple (target)	≤ 0.3 (base case)	—	Net Burn / Net New ARR; PLG products can be very efficient

3-year revenue model (base case)

Milestone	Timeline	ARR	Paying Operators
First 10 paying operators	Month 1 post-launch	~\$7K ARR	10
\$10K MRR	Month 3 post-launch (Sep 2026)	~\$120K ARR	~130 paying ops
\$50K MRR	Month 9 post-launch (Mar 2027)	~\$600K ARR	~650 paying ops
\$100K MRR	Month 14 post-launch (Aug 2027)	~\$1.2M ARR	~1,200 paying ops
\$500K MRR	Month 24 post-launch (Jun 2028)	~\$6M ARR	~5,000 paying ops
\$1M MRR	~Q3 2028	~\$12M ARR	~10,000+ paying ops

Path to \$1M ARR requires ~1,050 paid operators at blended \$80/month pricing — achievable with 2–3× the free user base and a 12–15% free-to-paid conversion rate within 30 days of signup.

Conservative scenario (P10): \$10K MRR by Nov 2026, cashflow-positive Apr–May 2027, end-2028 ARR ~\$2M. **Base case (P50):** \$10K MRR by Sep 2026, cashflow-positive Dec 2026, end-2028 ARR ~\$7.2M. **Optimistic (P90):** \$10K MRR by Aug 2026, cashflow-positive Sep 2026, end-2028 ARR ~\$25M.

Phase 2 revenue extension

When Gordon Desk is live and operator volume crosses the threshold that justifies venue partnerships, a transaction take-rate model becomes viable. At scale, a take-rate model has near-zero marginal cost of revenue, making the subscription tiers a floor, not a ceiling. Estimated rate: 0.05–0.15bps on notional execution volume routed through Gordon at scale.

At Personal Finance OS scale, the per-operator context (spending patterns, positions, income flows) can power embedded financial products (credit, insurance, tax) with revenue-share or embedded product fees. This is a 3–5 year horizon — do not optimize for it now, but do not build the data architecture in ways that make it impossible.

Growth model & go-to-market

PLG with community amplification

Gordon is a developer-adjacent, terminal-first product launching in the exact communities where PLG motions work — GitHub, HN, Finance Twitter, Claude Code community. SLG would require sales people, targeting buyers who are not yet in a discoverable pipeline, for a product that is best understood by doing. The decision is PLG.

The core insight on GTM: The Stuck Operator and the Crypto-Native Power User are concentrated in discoverable, high-signal communities. The goal in Phase 1 is not volume — it is signal quality from the right cohort.

Channel playbook

Phase 1 — Friends-alpha to public (now → Q3 2026):

Show HN	Canonical launch for developer-adjacent tools with technical story	Lead with architecture (plan-first preview, permission model) not "AI trading bot"
Finance Twitter / X	Highest concentration of Stuck Operators	"Vibe trading" vocabulary ownership; document alpha operator workflows publicly
GitHub public repo	CLI is the product; stars = credibility signal	Release with real README, real API docs, real demo
Claude Code / Cursor community	Highest-fit technical profile; understands agentic AI paradigm	Cross-post explicitly; "The Claude Code of trading" analogy
r/algotrading	Self-identifies as automation-intent Stuck Operators	Authentic feedback-seeking post format (see pitch-community.md)

Anti-channels (Phase 1): Paid acquisition, influencer sponsorships, and affiliate programs. Wrong for pre-public stage — priority is signal quality, not volume.

Phase 2 — Gordon Desk (12–18 months):

- YouTube + creator program: visual interface unlocks demo content
- Finance Twitter continued: own "vibe trading" vocabulary
- Partner/creator program: revenue share or free Power tier for creators who document workflows

Phase 3 — Personal Finance OS (24–36 months):

- Embedded distribution in partner products
- Bank/brokerage partnerships
- Different GTM motion — do not plan it now

First 100 paying operators

Cohort	Size	Channel	Why
Friends-alpha	20–30	Invited	Highest-trust; most feedback-rich

Finance Twitter + GitHub	30–40	Show HN + Twitter launch	Highest concentration of Stuck Operators
Crypto Discord	20–30	DeFi/quant Discord communities	Pre-existing fragmentation pain
Cursor/Claude Code crossover	10–15	Claude Code community, HN	Technical comfort + AI-first mindset

Key metrics & validation framework

North Star Metric

Operator-Weeks Active — unique operators who completed at least one full workflow cycle (scan → plan → preview → execute or paper-execute) in a given week.

This is the North Star because it measures the product doing its actual job: turning conviction into disciplined execution on a recurring basis. It is not gameable with a single click. It is a leading indicator of retention, data flywheel accumulation, and word-of-mouth.

One Metric That Matters (OMTM) — Phase 1

Plan-First Conversion Rate — percentage of operators who reach the preview step and confirm execution (rather than abandoning at the plan stage). Target: >60%.

This is the trust metric. If operators consistently abandon at the plan stage, the plan preview is generating anxiety rather than resolving it. If conversion is high, the trust mechanism is working and the safety architecture is an accelerant, not a barrier.

Core metrics table

Free → Pro conversion (30 days)	≥ 12%	Validates paper mode as an on-ramp; below 5% = structural onboarding failure
Pro → Power upgrade (6 months)	≥ 20%	Key expansion driver; Power tier has 3× ARPU and dramatically better LTV:CAC
Day 30 operator retention	≥ 25% (good); ≥ 35% (excellent)	Habit formation; leading indicator of long-term retention

Monthly churn (Pro)	$\leq 4.5\%$	Above 7%/month = structural PMF issue, not fixable by marketing
NDR at Month 12	$\geq 105\%$	Pro→Power upgrades drive expansion; target NDR > 100%
CAC (blended, Phase 1)	$\leq \$75$	PLG-dominant early growth should keep CAC low
Plans approved / operator / week	≥ 3	Core engagement signal; below 1/week = at-risk operator
Permission mode graduation (60 days)	$\geq 30\%$ paper→ask/auto	Validates trust ladder is working
Organic referral rate	$\geq 20\%$	Word-of-mouth is the primary growth mechanism at friends-alpha stage

The 3 metrics that make or break the model

1. Free → Pro conversion ($\geq 12\%$ in 30 days) — if paper trading doesn't create enough conviction to upgrade, the free tier is not working as an on-ramp.

2. Monthly churn ($\leq 4.5\%$ Pro, $\leq 2.5\%$ Power) — at \$49 ARPU, a 4.5% monthly churn gives ~22-month average lifetime. At 8% churn, LTV:CAC economics collapse.

3. Pro → Power upgrade ($\geq 20\%$ in 6 months) — the Power tier is 3× the revenue and has dramatically better economics. If operators are satisfied at Pro and never upgrade, ARPU stagnates.

Validation experiments (alpha cohort)

Experiment	Tests	Method
Plan-first abandonment heuristic	Whether trust mechanism converts	Log preview→confirm rate; exit prompt at abandon; 20 alpha ops is enough signal
Paper mode → live prompt at Day 14	Whether paper graduation is a real funnel	Single email after 3 paper cycles: "Your paper results are in. Ready to go live?"
Crypto-only vs. multi-venue cohort retention	Whether multi-venue span drives retention	Tag cohort at registration; compare Day 30 engagement and churn

The steelman — the strongest skeptic argument, answered

The strongest argument against us comes from Naeem Aslam ([Investing.com](https://www.investing.com)): "Performance is not validated in

days. It is revealed over cycles." AI trading systems that look good in stable conditions tend to fail during volatility, and an autonomous system executing real orders turns a flawed signal into a flawed action with direct financial consequences. Markets compress truth over years, not weeks. This is the serious version of the skepticism.

Three responses.

First, we agree with the premise and built the product around it. Gordon defaults to `ask` mode. `auto` mode exists but requires an explicit mandate and ships with risk gates, exposure caps, drawdown limits, and a kill switch. `paper` mode is a first-class citizen of the permission truth table, the safe on-ramp to live execution. We are not betting on agentic execution being validated in days. We are betting on a product that *lets the operator validate it over cycles* while keeping them in control of every cycle. That is the feature.

Second, the skeptic argument cuts deeper against the institutional side than against us. Institutional quant funds bet capital on models before out-of-sample cycles validate them. That is the whole point of trading. Our operators do not. Gordon's worst case in `ask` mode is that the operator rejects every trade and learns nothing. The institutional worst case is a fund blowing up. The system architecture shifts risk *away* from the operator, which is precisely the asymmetry retail investors currently suffer the other direction.

Third, the long-term case for agentic retail is stronger under this frame, not weaker. If performance is revealed over cycles, the product that wins is the one that keeps an operator in the game across cycles, preserving context, memory, discipline, and learning from mistakes. That is exactly the personal-context data flywheel Gordon is designed to build. Every competitor is optimized for today's trade. Gordon is optimized for an operator's decade of trades. Time is on our side.

Strategic self-assessment

Three honest perspectives on where General Liquidity stands today.

From a VC investor lens: Lean Invest. The architecture is genuinely differentiated — plan-first preview and six-mode permission model are architecture-level defensible and no competitor has both. Regulatory tailwinds are structural. Market timing is near-perfect. The constraint is that there is zero published activation data — investors are pricing on differentiation and timing alone. The LLM cost structure is a hidden unit economics variable that needs to be modeled against actual P90 usage data before public launch. **The category window is the clock: 6–12 months to establish position before a well-capitalized competitor invests in closing the gap.**

From a potential key hire lens: Lean Join. The technical problems are real and unusual — AI safety in financial execution, multi-venue integration engineering, the human-AI trust interface. The mission is coherent. The risk is pre-revenue equity discount combined with competitive timing pressure. Shifts to Strong Join if: public launch is within 60–90 days, alpha activation data is positive, and compensation reflects the stage risk honestly.

From the founder's lens — Health Grade: Stable. The most significant honest gaps: (1) activation data from the 20–30 alpha operators has not been formally extracted — the ICP validation status tracker explicitly marks "10+ interviews completed" as unvalidated; (2) no public presence means category vocabulary is being set by competitors right now; (3) an LLM execution error at this stage creates both user harm and a narrative event at the worst possible moment. The company is not in distress. It is at the inflection point where the next 60–90 days determine whether "stable" becomes "strong."

Cross-perspective bottom line: All three lenses agree on the same thing: the plan-first preview is the trust mechanism the category is missing, and what happens with the alpha cohort in the next 30–60 days determines the company's trajectory more than any other factor.

Top 3 priority actions (next 90 days):

1. Extract activation evidence from alpha operators — define the activation metric precisely (connected venue + plan-first preview completed + paper trade executed) and measure it for all 20–30 alpha operators now.
2. Stake the category claim publicly before full product launch — publish a Show HN, a "vibe trading manifesto," or a technical writeup on the plan-first architecture. Claim the vocabulary before a competitor does.
3. Document the plan-first preview architecture as IP — the window to establish earliest priority date is open now, pre-public, pre-copy. A provisional filing establishes priority and signals intentional architecture to future acquirers and investors.

10-year vision

In 2036, General Liquidity is the financial operating system for the agentic generation. Gordon, in its terminal form, its Desk form, and its personal-finance form, is how a material fraction of Gen Z and younger millennials actually operate their money.

- Tens of millions of operators have accumulated multi-year financial context inside Gordon. That context is the primary asset of the company: every plan, every trade, every mistake, every calibration made, portable and compounding.
- Gordon Desk is the default environment for any serious retail operator who wants to run strategies, backtests, and cross-venue portfolios. It is to personal trading what Cursor is to coding — the place where humans and agents collaborate at the level of an agentic Bloomberg Terminal, with shared context and real authority. Agentic multiplayer is a normal work pattern. Teams of two, three, or ten operators run shared theses together with shared agent infrastructure.
- Gordon agents transact, pay, negotiate, reconcile, and optimize across cards, banks, wallets, exchanges, payroll, and subscriptions. The boundary between trading and personal finance has dissolved because both are just forms of delegated economic action. Every external merchant, insurer, bank, and broker builds API-first, machine-legible products because Gordon, not a human, is increasingly the counterparty.

- The applied-research arm has set the vocabulary of the control-plane category in public, on the record, and has become a reference in how trading-native AI safety is defined and measured.
- The company is the canonical worldbuilding answer to the question "*what does it mean for ordinary people to have economic agency in the age of agents?*" The answer is visible, operable, profitable, and used.

If we are right about programmability without agency being the trap of the next decade, General Liquidity is the company that broke the trap for everyone who was not already on the right side of it.

Supporting materials

The following reference documents are available in `docs/pitch/general-liquidity/`:

<code>positioning-doc.md</code>	Full April Dunford 5+1 positioning analysis
<code>positioning-statement.md</code>	Moore statement, Neumeier onliness test, tagline candidates, elevator pitch, one-liner variants
<code>icp.md</code>	Full ICP definition, persona canvases, buying center map, pain hierarchy scoring, interview questions
<code>competitors-report.md</code>	Competitive landscape executive summary, market concentration, strategic opportunities and risks
<code>competitive-matrix.md</code>	Feature comparison table: 14 features × 8 competitors
<code>pricing-landscape.md</code>	Tier-by-tier pricing comparison, value metric analysis, pricing whitespace
<code>battle-cards/public-com.md</code>	Battle card: Public.com
<code>battle-cards/3commas-quantpilot.md</code>	Battle card: 3Commas QuantPilot
<code>battle-cards/tradingview-manual.md</code>	Battle card: TradingView + manual execution
<code>battle-cards/status-quo.md</code>	Battle card: inertia / status quo
<code>startup-canvas.md</code>	Product strategy canvas: 9-section strategy + business model, coherence check, hypotheses, experiments
<code>financial-model.md</code>	3-year revenue model (3 scenarios), unit economics, milestones, sensitivity analysis

metrics-framework.md	North Star, OMTM, full metrics suite, investor dashboard format, validation experiments
startup-analysis.md	Three-perspective stress test: VC investor, key hire, CEO/founder
pitch-operator.md	Operator pitch: 30-sec, 2-min, Product Hunt, App Store/GitHub formats + objection responses
pitch-community.md	Community pitch: HN Show HN, Twitter/X thread, Reddit, Discord
pitch-partner.md	Partner pitch: venue partners, API partners, integration brief
pitch-appendix.md	Q&A prep: 10 operator objections, 5 press questions, 5 partner objections, pitch scorecard

© 2026 General Liquidity, Inc.